

Qiang Xu

AI Product Management Intern

LLM Applications | Agents | Product Evaluation

Beijing, China | +86 15838695058 | 3420260662@bit.edu.cn

Portfolio: <https://steve-ron.github.io/ai-product-ops-portfolio/>

M.S. student at BIT | Weekend classes | Available 5 weekdays/week | 6+ months

PROFILE

Graduate student focused on **LLM applications, agents, and product evaluation**. Translates ambiguous model-experience and user-feedback problems into requirements, PRDs, evaluation frameworks, traceable datasets, and interactive prototypes, with explicit acceptance criteria and iteration hypotheses.

2 interactive MVPs | 2 PRDs | 48 evaluation cases | 120 preference records | 156 user-feedback records

Portfolio evidence: 5 responsive pages, 2 PRDs, public-source datasets, and automated validation scripts; works online and offline, with every key count traceable to source files.

EDUCATION

Beijing Institute of Technology | M.S. in Digital Economy (AI & Economic Decision-Making) · Part-time Sep 2026 - Jun 2028

Chengdu College of Arts and Sciences | B.A. in Calligraphy · Full-time Sep 2021 - Jun 2025

SELECTED AI PRODUCT PROJECTS

LLM Lens | LLM Evaluation & Bad-Case Workspace Aug 2026

Independent project | Product Lead (AI-assisted implementation) | PRD · Prototype · Data / Evaluation

- Product definition:** Framed the problem of inconsistent model-experience reviews; produced the PRD and interactive MVP, modeling tasks, raw outputs, scores, root-cause hypotheses, review status, and acceptance criteria.
- Evaluation system:** Built an **8-dimension** framework covering instruction following, multi-turn context, factuality and recency, safety boundaries, and agent tool awareness; designed **48 standardized cases**, all transparently marked pending execution.
- MVP and data:** Deterministically sampled **120** bilingual LMArena preference records while retaining prompts, full responses, model versions, winner labels, and source links; implemented combined filters, response comparison, local scoring/version records, and JSON/CSV export.
- Product judgment:** Separated public human preferences, rule-based screening, and human conclusions; scoped the next loop as dual-rater agreement -> live API baseline -> cost, latency, and tool-trace analysis.

Signal Desk | AI Product User Feedback Insight Workspace Aug 2026

Independent project | Product Lead (AI-assisted implementation) | PRD · Prototype · User Insight / Analytics

- Product definition:** Addressed fragmented feedback and the gap between frequency and priority; defined a traceable feedback ledger, taxonomy, explainable Triage Score, human-review workflow, and interactive MVP in the PRD.
- Data governance:** Used the public GitHub API to collect Open WebUI Issues; removed pull requests, usernames, and bodies, producing **156 traceable records** containing only public metadata and source links.
- User insight:** Built **9 themes** and **3 feedback types**; identified issue clusters for model/API connectivity (**45**), UI/conversation flow (**22**), and knowledge base/RAG (**21**), then proposed connection diagnostics, visible RAG stages, and manual review of the Other category.
- MVP and boundaries:** Combined comments, reactions, status, and blocking signals into an explainable Triage Score with multi-filter search, source-issue traceability, and CSV export; explicitly positioned the score as a review queue, not a true business priority.

CORE SKILLS

Product: problem framing / requirements / PRDs / workflow design / prototyping / acceptance criteria / user-feedback analysis / product evaluation / bad-case analysis

AI: LLM applications and limitations / agents / prompting / RAG fundamentals / test-set design / evaluation metrics

Data & tools: Excel (PivotTables/functions) / Python and SQL fundamentals / GitHub API / JSON and CSV / AI-assisted prototyping

English: CET-4 516; CET-6 475; able to read technical documentation and public dataset specifications